

# This Assignment was done by :

Mahmoud Labib Elsamadony\_\_\_\_20208029

Ahmed Alaa Elsaadani\_\_\_\_20208003

Abdelrahman Mezar\_\_\_\_20208018

Mohamed Amr Shehab\_\_\_\_20208041

Mohamed Ahmed Seif\_\_\_\_20208025

## RSNA Pneumonia Detection

1- The dataset provided in two folder:

- train\_images.zip
- test\_images.zip.

2- train.csv:

- the training set. Contains patientIds and bounding box / target information.
- patientId :- A patientId. Each patientId corresponds to a unique image.
- x :- the upper-left x coordinate of the bounding box.
- y :- the upper-left y coordinate of the bounding box.
- width :-the width of the bounding box.
- height :- the height of the bounding box.
- Target :- the binary Target, indicating whether this sample has evidence of pneumonia

## First we view the data

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import pydicom
from PIL import Image
import os
from ultralytics import YOLO
```

```
In [ ]: df = pd.read_csv('dataset/train_labels.csv')
df.head()
```

Out [ ]:

	patientId	x	y	width	height	Target
0	0004cfab-14fd-4e49-80ba-63a80b6bddd6	NaN	NaN	NaN	NaN	0
1	00313ee0-9eaa-42f4-b0ab-c148ed3241cd	NaN	NaN	NaN	NaN	0
2	00322d4d-1c29-4943-afc9-b6754be640eb	NaN	NaN	NaN	NaN	0
3	003d8fa0-6bf1-40ed-b54c-ac657f8495c5	NaN	NaN	NaN	NaN	0
4	00436515-870c-4b36-a041-de91049b9ab4	264.0	152.0	213.0	379.0	1

In [ ]:

```
# replace nans with 0
df.fillna(0, inplace=True)
df.head()
```

Out [ ]:

	patientId	x	y	width	height	Target
0	0004cfab-14fd-4e49-80ba-63a80b6bddd6	0.0	0.0	0.0	0.0	0
1	00313ee0-9eaa-42f4-b0ab-c148ed3241cd	0.0	0.0	0.0	0.0	0
2	00322d4d-1c29-4943-afc9-b6754be640eb	0.0	0.0	0.0	0.0	0
3	003d8fa0-6bf1-40ed-b54c-ac657f8495c5	0.0	0.0	0.0	0.0	0
4	00436515-870c-4b36-a041-de91049b9ab4	264.0	152.0	213.0	379.0	1

In [ ]:

```
df.describe()
```

Out [ ]:

	x	y	width	height	Target
count	30227.000000	30227.000000	30227.000000	30227.000000	30227.000000
mean	124.561683	115.960962	69.060575	104.084825	0.316108
std	216.326397	190.012883	106.910496	176.932152	0.464963
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	193.000000	231.000000	169.000000	188.000000	1.000000
max	835.000000	881.000000	528.000000	942.000000	1.000000

In [ ]:


```
# Plotting the first image

path_to_image = df['patientId'][0]

# Load the .dcm image
ds = pydicom.dcmread(f'dataset/train_images/{path_to_image}.dcm')

# Access image data
image = ds.pixel_array

# Display the image
plt.imshow(image, cmap='gray')
plt.show()
```



```
In [ ]: # Plotting the first image with bounding box

# Load the .dcm image
path_to_image = df['patientId'][4]
ds = pydicom.dcmread(f'dataset/train_images/{path_to_image}.dcm')

# Access image data
image = ds.pixel_array

# Display the image
plt.imshow(image, cmap='gray')

# Add bounding box


# Get the target values
x = df['x'][4]
y = df['y'][4]
width = df['width'][4]
height = df['height'][4]

# Create a Rectangle patch

rect = plt.Rectangle((x, y), width, height, edgecolor='r', facecolor='none')

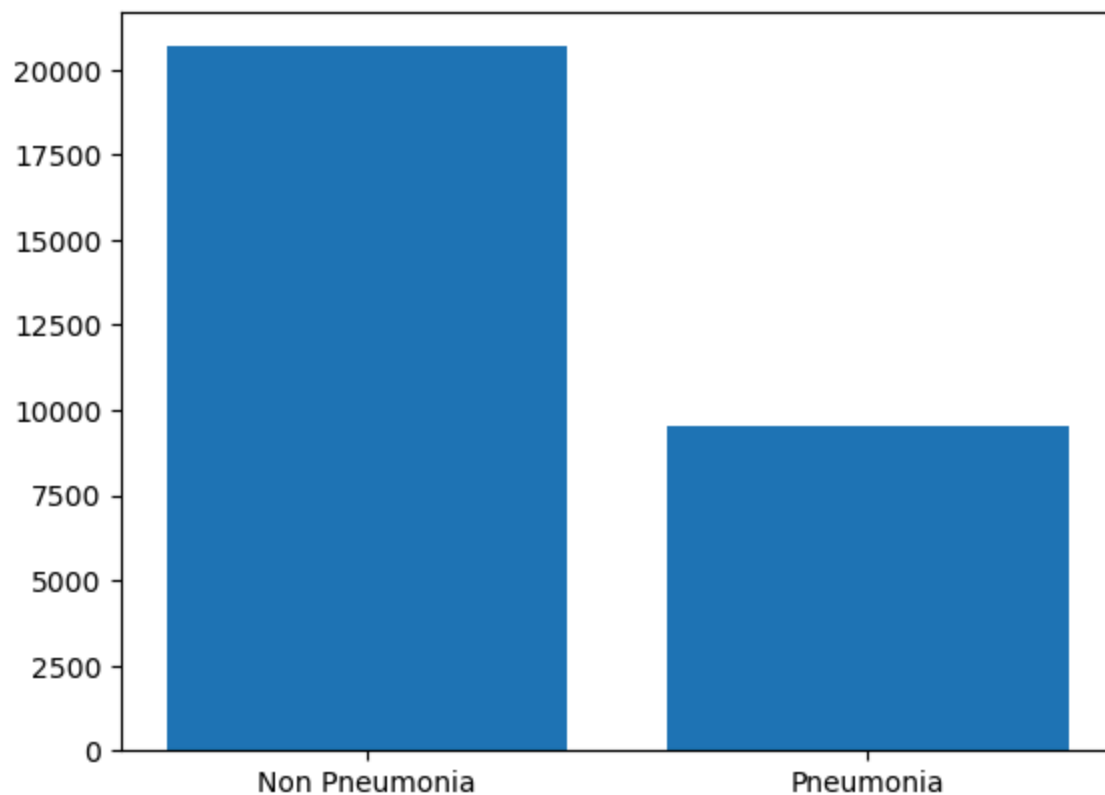
# Add the patch to the Axes
plt.gca().add_patch(rect)

plt.show()
```



```
In [ ]: # compare between the number of images with pneumonia and without pneumonia
print(df['Target'].value_counts())
plt.bar(['Non Pneumonia', 'Pneumonia'], df['Target'].value_counts());
```

```
Target
0    20672
1     9555
Name: count, dtype: int64
```



```
In [ ]: # count the .dcm files in the dataset/train_images folder
print(len(os.listdir('dataset/train_images')))
```

26684

```
In [ ]: # check if any of the patientId in the train_labels.csv file are in the dataset/test_images
test_images = os.listdir('dataset/test_images')

for patientId in df['patientId']:
    if f'{patientId}.dcm' in test_images:
        print(patientId)
```

```
In [ ]: df.shape
```

```
Out[ ]: (30227, 6)
```

```
In [ ]: # count the number of unique patients
df['patientId'].nunique()
```

```
Out[ ]: 26684
```

```
In [ ]: print(len(os.listdir('dataset/test_images')))
```

3000

- Nan values are replaced with zeros
- number of images with non-pneumonia are double the images with pneumonia (imbalanced data)
- number of unique patients in the dataframe is equal to the number of training images (no redundancy)

## Convert the .dcm images to a format that YOLOv8 can work with

- First, we extract the image data from .dcm file
- Second, we will Normalize the pixel values to be from 0 - 255 (leading to faster convergence)
- we then convert it to jpg format

```
In [ ]: # convert the .dcm files to .jpg files

def convert_dcm_to_jpg(input_folder, output_folder):
    for file in os.listdir(input_folder):
        if os.path.isfile(f'{output_folder}/{file.split(".")[0]}.jpg'):
            continue
        ds = pydicom.dcmread(f'{input_folder}/{file}')
        img = ds.pixel_array
        # Normalize the pixel values to the range [0, 255]
        image_normalized = (img - np.min(img)) / (np.max(img) - np.min(img)) * 255
        image_normalized = image_normalized.astype(np.uint8) # Convert to unsigned 8-bit

        # Convert the numpy array to a PIL image
        image = Image.fromarray(image_normalized)

        filename = file.split('.')[0]
        image.save(f'{output_folder}/{filename}.jpg')
```

```
In [ ]: convert_dcm_to_jpg('dataset/train_images', 'dataset/images/train')
```

```
In [ ]: convert_dcm_to_jpg('dataset/test_images', 'dataset/images/test')
```

# Normalizing the labels (YOLO Format)

```
In [ ]: def normalize_bboxes(df):
        img_size = 1024

        # normalize the bounding boxes
        df['width'] = df['width'] / img_size
        df['height'] = df['height'] / img_size
        # Centering X , Y coordinates and normalize them
        df['x'] = df['x'] / img_size + df['width'] / 2
        df['y'] = df['y'] / img_size + df['height'] / 2

def denormalize_bbox(rx, ry, rw, rh):
    img_size = 1024

    x = (rx-rw/2)*img_size
    y = (ry-rh/2)*img_size
    w = rw*img_size
    h = rh*img_size

    return x, y, w, h
```

```
In [ ]: normalize_bboxes(df)

df.head()
```

```
Out[ ]:
```

	patientId	x	y	width	height	Target
0	0004cfab-14fd-4e49-80ba-63a80b6bddd6	0.000000	0.000000	0.000000	0.000000	0
1	00313ee0-9eaa-42f4-b0ab-c148ed3241cd	0.000000	0.000000	0.000000	0.000000	0
2	00322d4d-1c29-4943-afc9-b6754be640eb	0.000000	0.000000	0.000000	0.000000	0
3	003d8fa0-6bf1-40ed-b54c-ac657f8495c5	0.000000	0.000000	0.000000	0.000000	0
4	00436515-870c-4b36-a041-de91049b9ab4	0.361816	0.333496	0.208008	0.370117	1

## Preparing The labels folder

```
In [ ]: # Converting the df to .txt files per patientId each patientId in separate file (as ther

def convert_df_to_txt_files(df, output_folder):
    if not os.path.exists(output_folder):
        os.makedirs(output_folder)
    for patientId in df['patientId'].unique():
        patient_df = df[df['patientId'] == patientId].copy()
        if patient_df['Target'].values[0] == 0:
            # if the patient doesn't have pneumonia, output empty txt file
            with open(f'{output_folder}/{patientId}.txt', 'w') as f:
                f.close()
            continue
        patient_df.drop(columns='patientId', inplace=True)
        # getting the target column to the first column
        cols = patient_df.columns.tolist()
        cols = cols[-1:] + cols[:-1]
        patient_df = patient_df[cols]
```

```
patient_df['Target'] = 0
patient_df.to_csv(f'{output_folder}/{patientId}.txt', index=False, header=False,
```

```
In [ ]: convert_df_to_txt_files(df, 'dataset/labels/train')
```

## Splitting the data into train\_1 , train\_2 , train\_3 and validation (Hardware constraints)

```
In [ ]: ## split the data into train_1 , train_2 , train_3 and validation sets and without shuffle

# get the unique patientIds
patientIds = df['patientId'].unique()

train_1 = patientIds[:9000]
train_2 = patientIds[9000:18000]
train_3 = patientIds[18000:24000]
validation = patientIds[24000:]

# get the patientIds for each set
train_1_df = df[df['patientId'].isin(train_1)].copy()
train_2_df = df[df['patientId'].isin(train_2)].copy()
train_3_df = df[df['patientId'].isin(train_3)].copy()
validation_df = df[df['patientId'].isin(validation)].copy()

# convert the train_1_df, train_2_df, train_3_df and validation_df to .txt files
# convert_df_to_txt_files(train_1_df, 'dataset/labels/train_1')
# convert_df_to_txt_files(train_2_df, 'dataset/labels/train_2')
# convert_df_to_txt_files(train_3_df, 'dataset/labels/train_3')
# convert_df_to_txt_files(validation_df, 'dataset/labels/validation')

# copy the images to the respective folders
def copy_images_to_folder(df, input_folder, output_folder):
    if not os.path.exists(output_folder):
        os.makedirs(output_folder)
    for patientId in df['patientId'].unique():
        os.system(f'cp {input_folder}/{patientId}.jpg {output_folder}/{patientId}.jpg')

# copy_images_to_folder(train_1_df, 'dataset/images/train', 'dataset/images/train_1')
# copy_images_to_folder(train_2_df, 'dataset/images/train', 'dataset/images/train_2')
# copy_images_to_folder(train_3_df, 'dataset/images/train', 'dataset/images/train_3')
# copy_images_to_folder(validation_df, 'dataset/images/train', 'dataset/images/validation')
```

## Training The YOLOv8 Algorithm

- we ran each part (train\_1 , train\_2 , train\_3) of the data for two epochs to achieve better results
- at the end of each part we get the best.pt from the previous one for the next training session
- that lead to better losses , precision and recall

```
In [ ]: model = YOLO("dataset/best3.pt")

res = model.train(data="data.yaml", epochs=1) # one final epoch to finish training
```

New <https://pypi.org/project/ultralytics/8.2.19> available 😊 Update with 'pip install -U ultralytics'  
Ultralytics YOLOv8.2.18 🚀 Python-3.11.5 torch-2.2.2 CUDA:0 (NVIDIA GeForce GTX 1650, 390 4MiB)

**engine/trainer:** task=detect, mode=train, model=dataset/best3.pt, data=data.yaml, epochs=1, time=None, patience=100, batch=16, imgsz=640, save=True, save\_period=-1, cache=False, device=None, workers=8, project=None, name=train, exist\_ok=False, pretrained=True, optimizer=auto, verbose=True, seed=0, deterministic=True, single\_cls=False, rect=False, cos\_lr=False, close\_mosaic=10, resume=False, amp=True, fraction=1.0, profile=False, freeze=None, multi\_scale=False, overlap\_mask=True, mask\_ratio=4, dropout=0.0, val=True, split=val, save\_json=False, save\_hybrid=False, conf=None, iou=0.7, max\_det=300, half=False, dnn=False, plots=True, source=None, vid\_stride=1, stream\_buffer=False, visualize=False, augment=False, agnostic\_nms=False, classes=None, retina\_masks=False, embed=None, show=False, save\_frames=False, save\_txt=False, save\_conf=False, save\_crop=False, show\_labels=True, show\_conf=True, show\_boxes=True, line\_width=None, format=torchscript, keras=False, optimize=False, int8=False, dynamic=False, simplify=False, opset=None, workspace=4, nms=False, lr0=0.01, lrf=0.01, momentum=0.937, weight\_decay=0.0005, warmup\_epochs=3.0, warmup\_momentum=0.8, warmup\_bias\_lr=0.1, box=7.5, cls=0.5, dfl=1.5, pose=12.0, kobj=1.0, label\_smoothing=0.0, nb\_s=64, hsv\_h=0.015, hsv\_s=0.7, hsv\_v=0.4, degrees=0.0, translate=0.1, scale=0.5, shear=0.0, perspective=0.0, flipud=0.0,fliplr=0.5, bgr=0.0, mosaic=1.0, mixup=0.0, copy\_paste=0.0, auto\_augment=randaugument, erasing=0.4, crop\_fraction=1.0, cfg=None, tracker=botsort.yaml, save\_dir=runs/detect/train

	from	n	params	module	arguments
0	-1	1	464	ultralytics.nn.modules.conv.Conv	[3, 1
6, 3, 2]					
1	-1	1	4672	ultralytics.nn.modules.conv.Conv	[16, 3
2, 3, 2]					
2	-1	1	7360	ultralytics.nn.modules.block.C2f	[32, 3
2, 1, True]					
3	-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 6
4, 3, 2]					
4	-1	2	49664	ultralytics.nn.modules.block.C2f	[64, 6
4, 2, True]					
5	-1	1	73984	ultralytics.nn.modules.conv.Conv	[64, 1
28, 3, 2]					
6	-1	2	197632	ultralytics.nn.modules.block.C2f	[128,
128, 2, True]					
7	-1	1	295424	ultralytics.nn.modules.conv.Conv	[128,
256, 3, 2]					
8	-1	1	460288	ultralytics.nn.modules.block.C2f	[256,
256, 1, True]					
9	-1	1	164608	ultralytics.nn.modules.block.SPPF	[256,
256, 5]					
10	-1	1	0	torch.nn.modules.upsampling.Upsample	[None,
2, 'nearest']					
11	[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
12	-1	1	148224	ultralytics.nn.modules.block.C2f	[384,
128, 1]					
13	-1	1	0	torch.nn.modules.upsampling.Upsample	[None,
2, 'nearest']					
14	[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
15	-1	1	37248	ultralytics.nn.modules.block.C2f	[192,
64, 1]					
16	-1	1	36992	ultralytics.nn.modules.conv.Conv	[64, 6
4, 3, 2]					
17	[-1, 12]	1	0	ultralytics.nn.modules.conv.Concat	[1]
18	-1	1	123648	ultralytics.nn.modules.block.C2f	[192,
128, 1]					
19	-1	1	147712	ultralytics.nn.modules.conv.Conv	[128,
128, 3, 2]					



```

20          [-1, 9] 1      0      ultralytics.nn.modules.conv.Concat      [1]
21          -1 1      493056 ultralytics.nn.modules.block.C2f      [384,
256, 1]
22      [15, 18, 21] 1      751507 ultralytics.nn.modules.head.Detect      [1, [6
4, 128, 256]]
Model summary: 225 layers, 3011043 parameters, 3011027 gradients, 8.2 GFLOPs


```

Transferred 355/355 items from pretrained weights

**TensorBoard:** Start with 'tensorboard --logdir runs/detect/train', view at <http://localhost:6006/>

Freezing layer 'model.22.dfl.conv.weight'

**AMP:** running Automatic Mixed Precision (AMP) checks with YOLOv8n...

**AMP:** checks passed 

**train:** Scanning /media/mahmoud/31A5D0D7034B3BD5/General/FCAI/8th Term/Assignment\_1/dataset/labels/train\_1.cache... 9000 images, 6932 backgrounds, 0 corrupt: 100%|██████████| 9000/9000 [00:00<?, ?it/s]

**val:** Scanning /media/mahmoud/31A5D0D7034B3BD5/General/FCAI/8th Term/Assignment\_1/dataset/labels/validation.cache... 2684 images, 2079 backgrounds, 0 corrupt: 100%|██████████| 2684/2684 [00:00<?, ?it/s]

Plotting labels to runs/detect/train/labels.jpg...

**optimizer:** 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automatically...

**optimizer:** AdamW(lr=0.002, momentum=0.9) with parameter groups 57 weight(decay=0.0), 64 weight(decay=0.0005), 63 bias(decay=0.0)

**TensorBoard:** model graph visualization added 

Image sizes 640 train, 640 val

Using 8 dataloader workers

Logging results to **runs/detect/train**

Starting training for 1 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
1/1	2.31G	2.032	2.694	1.866	1	640: 100% ██████████		
██████████  563/563 [08:51<00:00, 1.06it/s]								
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95):	
100% ██████████	84/84	[00:59<00:00, 1.42it/s]						
	all	2684	981	0.153	0.243	0.108	0.0331	

1 epochs completed in 0.167 hours.

Optimizer stripped from runs/detect/train/weights/last.pt, 6.2MB

Optimizer stripped from runs/detect/train/weights/best.pt, 6.2MB

Validating runs/detect/train/weights/best.pt...

Ultralytics YOLOv8.2.18 🚀 Python-3.11.5 torch-2.2.2 CUDA:0 (NVIDIA GeForce GTX 1650, 3904MiB)

Model summary (fused): 168 layers, 3005843 parameters, 0 gradients, 8.1 GFLOPs

	Class	Images	Instances	Box(P	R	mAP50	mAP50-95):	
100% ██████████	84/84	[00:51<00:00, 1.62it/s]						
	all	2684	981	0.157	0.224	0.109	0.0337	

Speed: 0.4ms preprocess, 17.1ms inference, 0.0ms loss, 0.9ms postprocess per image

Results saved to **runs/detect/train**

## Predicting Bounding Boxes

```

In [ ]: model = YOLO("best3.pt")
        results = model('dataset/images/test', conf=0.1, stream=True)

def get_id_from_path(path):
    return os.path.splitext(os.path.basename(path))[0]

```

```
for result in results:
    id = get_id_from_path(result.path)
    with open(f'dataset/predictions/{id}.txt', 'w') as f:
        if len(result.xywh) == 0:
            f.write(f'{id}\n')
        else:
            with open(f'dataset/predictions/{id}.txt', 'w') as f:

                x, y, w, h = denormalize_bbox(*result.xywh.tolist())
                confidence = result.conf
                f.write(f'{id} {confidence} {x} {y} {w} {h} \n')
```